

Variational Auto Encoder - Mnist

https://github.com/keras-team/keras/blob/master/examples/variational_autoencoder.py

because this one is broken: <https://blog.keras.io/building-autoencoders-in-keras.html>

```
[1] from keras.datasets import mnist
import numpy as np
```

Using TensorFlow backend.

```
[2] from keras.layers import Input, Lambda, Dense
from keras.models import Model
```

```
[3] from keras import backend as K
```

```
[20] # network parameters
original_dim=784
input_shape = (original_dim, )
intermediate_dim = 512
batch_size = 128
latent_dim = 2
epochs = 50
```

```
[67] inputs = Input(shape=input_shape, name='encoder_input')
x = Dense(intermediate_dim, activation='relu')(inputs)

inputs.shape, x.shape
```

```
(TensorShape([None, 784]), TensorShape([None, 512]))
```

```
[68] (x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = x_train.reshape((len(x_train),
np.prod(x_train.shape[1:])))
```

```
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))
```

```
[69] z_mean = Dense(latent_dim, name='z_mean')(x)
      z_log_var = Dense(latent_dim, name='z_log_var')(x)
```

```
[70] def sampling(args):
      z_mean, z_log_var = args
      batch = K.shape(z_mean)[0]
      dim = K.int_shape(z_mean)[1]
      # by default, random_normal has mean = 0 and std = 1.0
      epsilon = K.random_normal(shape=(batch, dim))
      return z_mean + K.exp(0.5 * z_log_var) * epsilon
```

```
[71] z = Lambda(sampling, output_shape=(latent_dim,), name='z')
      ([z_mean, z_log_var])
```

```
[83] # instantiate encoder model
      encoder = Model(inputs, [z_mean, z_log_var, z], name='encoder')
      encoder.summary()
      plot_model(encoder, to_file='vae_mlp_encoder.png',
      show_shapes=True)
```

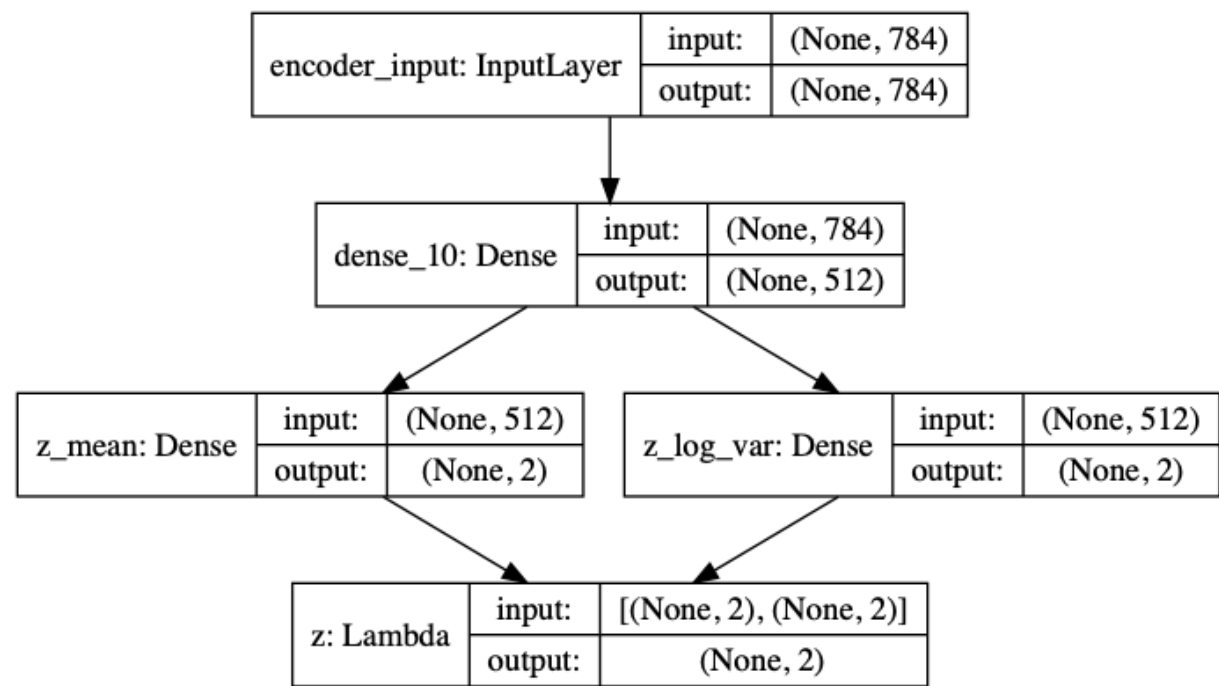
Model: "encoder"

Layer (type)	Output Shape	Param #
encoder_input (InputLayer)	(None, 784)	0
dense_10 (Dense)	(None, 512)	401920
z_mean (Dense)	(None, 2)	1026
z_log_var (Dense)	(None, 2)	1026

```
z (Lambda) (None, 2) 0
z_mean[0][0]
```

```
z_log_var[0][0]
```

```
=====
Total params: 403,972
Trainable params: 403,972
Non-trainable params: 0
-----
```



```
[84] # build decoder model
latent_inputs = Input(shape=(latent_dim,), name='z_sampling')
x = Dense(intermediate_dim, activation='relu')(latent_inputs)
outputs = Dense(original_dim, activation='sigmoid')(x)
```

```
[85] # instantiate decoder model
decoder = Model(latent_inputs, outputs, name='decoder')
decoder.summary()
plot_model(decoder, to_file='vae_mlp_decoder.png',
show_shapes=True)
```

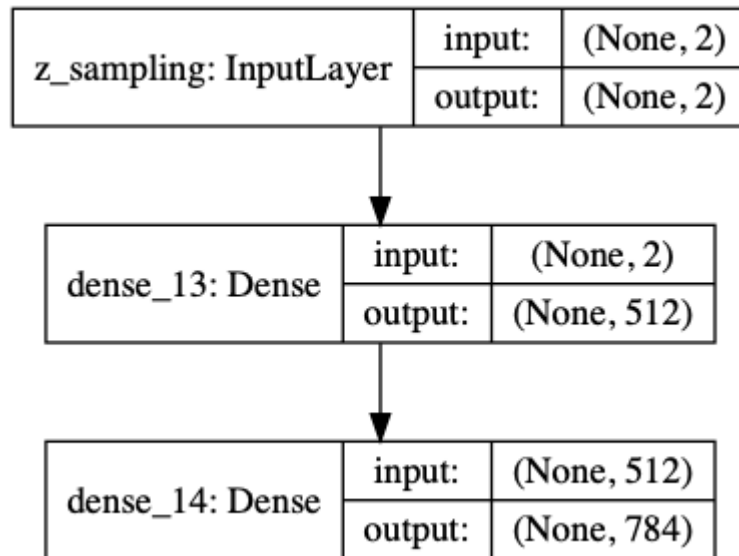
```
Model: "decoder"
```

```
-----
Layer (type)                Output Shape              Param #
-----
z_sampling (InputLayer)     (None, 2)                 0
-----
```

dense_13 (Dense)	(None, 512)	1536

dense_14 (Dense)	(None, 784)	402192
=====		

Total params: 403,728
Trainable params: 403,728
Non-trainable params: 0



```
[89] # end-to-end autoencoder
#vae = Model(x, x_decoded_mean)
# instantiate VAE model
outputs = decoder(encoder(inputs)[2])
vae = Model(inputs, outputs, name='vae_mlp')
```

```
[90] # generator, from latent space to reconstructed inputs
decoder_input = Input(shape=(latent_dim,))
_h_decoded = decoder_h(decoder_input)
_x_decoded_mean = decoder_mean(_h_decoded)
generator = Model(decoder_input, _x_decoded_mean)
```

```
[91] from keras.losses import mse, binary_crossentropy
```

```
[92] # train
def vae_loss(x, x_decoded_mean):
    #xent_loss = objectives.binary_crossentropy(x,
    x_decoded_mean)
    inputs, outputs = x, x_decoded_mean
    reconstruction_loss = mse(inputs, outputs)
    reconstruction_loss *= original_dim
    kl_loss = 1 + z_log_var - K.square(z_mean) - K.exp(z_log_var)
```

```

kl_loss = K.sum(kl_loss, axis=-1)
kl_loss *= -0.5
return K.mean(reconstruction_loss + kl_loss)

```

```
[93] vae.compile(optimizer='adam', loss=vae_loss)
```

```
[94] vae.summary()
```

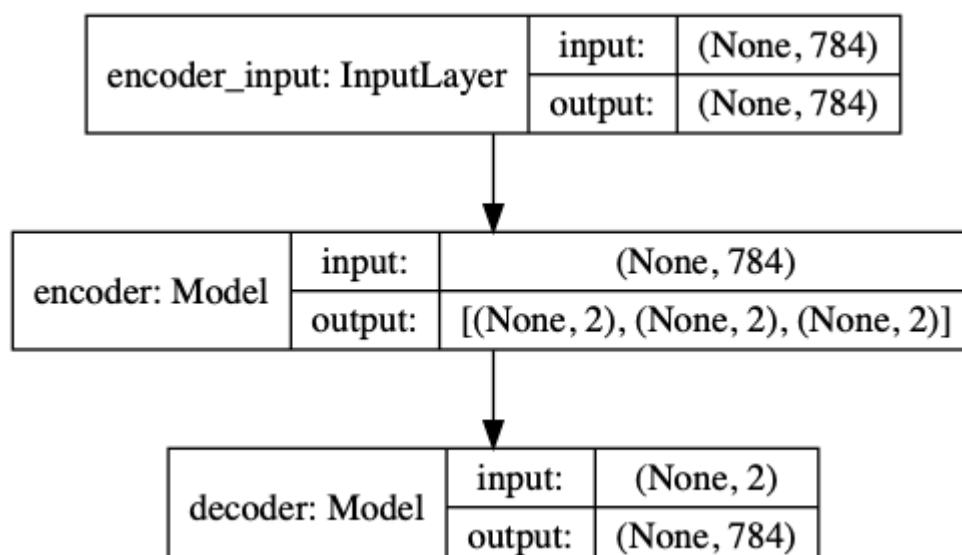
Model: "vae_mlp"

Layer (type)	Output Shape	Param #
encoder_input (InputLayer)	(None, 784)	0
encoder (Model)	[(None, 2), (None, 2), (N 403972	
decoder (Model)	(None, 784)	403728

=====
 Total params: 807,700
 Trainable params: 807,700
 Non-trainable params: 0
 =====

```
[95] from keras.utils import plot_model
```

```
[96] plot_model(vae,
               to_file='vae_mlp.png',
               show_shapes=True)
```



```
[97] history = vae.fit(x_train, x_train,
                      shuffle=True,
                      epochs=epochs,
                      batch_size=batch_size,
                      validation_data=(x_test, x_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/50

60000/60000 [=====] - 8s 134us/step - loss: 52.9492 - val_loss: 44.7417

Epoch 2/50

60000/60000 [=====] - 8s 136us/step - loss: 43.9253 - val_loss: 43.0114

Epoch 3/50

60000/60000 [=====] - 9s 154us/step - loss: 42.7439 - val_loss: 42.3133

Epoch 4/50

60000/60000 [=====] - 9s 151us/step - loss: 42.0549 - val_loss: 41.6784

Epoch 5/50

60000/60000 [=====] - 10s 171us/step - loss: 41.4746 - val_loss: 41.1764

Epoch 6/50

60000/60000 [=====] - 9s 149us/step - loss: 41.0069 - val_loss: 40.7550

Epoch 7/50

60000/60000 [=====] - 9s 144us/step - loss: 40.5779 - val_loss: 40.4166

Epoch 8/50

60000/60000 [=====] - 9s 145us/step - loss: 40.2560 - val_loss: 40.3243

Epoch 9/50

60000/60000 [=====] - 8s 141us/step - loss: 39.9779 - val_loss: 39.9516

Epoch 10/50

60000/60000 [=====] - 9s 144us/step - loss: 39.7271 - val_loss: 39.6259

Epoch 11/50

60000/60000 [=====] - 9s 148us/step - loss: 39.4854 - val_loss: 39.6053

Epoch 12/50

60000/60000 [=====] - 10s 171us/step - loss: 39.2840 - val_loss: 39.4055

Epoch 13/50

60000/60000 [=====] - 9s 154us/step - loss: 39.1240 - val_loss: 39.2507

Epoch 14/50

60000/60000 [=====] - 9s 158us/step - loss: 38.9606 - val_loss: 39.1248

Epoch 15/50

60000/60000 [=====] - 8s 140us/step - loss: 38.8271 - val_loss: 38.9282

Epoch 16/50
60000/60000 [=====] - 9s 147us/step - loss:
38.6658 - val_loss: 38.8016
Epoch 17/50
60000/60000 [=====] - 8s 139us/step - loss:
38.5476 - val_loss: 38.8427
Epoch 18/50
60000/60000 [=====] - 9s 149us/step - loss:
38.4409 - val_loss: 38.6343
Epoch 19/50
60000/60000 [=====] - 9s 149us/step - loss:
38.3461 - val_loss: 38.5794
Epoch 20/50
60000/60000 [=====] - 9s 147us/step - loss:
38.2301 - val_loss: 38.4719
Epoch 21/50
60000/60000 [=====] - 8s 138us/step - loss:
38.1407 - val_loss: 38.5420
Epoch 22/50
60000/60000 [=====] - 9s 151us/step - loss:
38.0807 - val_loss: 38.4466
Epoch 23/50
60000/60000 [=====] - 8s 138us/step - loss:
38.0109 - val_loss: 38.3400
Epoch 24/50
60000/60000 [=====] - 8s 137us/step - loss:
37.9214 - val_loss: 38.2518
Epoch 25/50
60000/60000 [=====] - 8s 138us/step - loss:
37.8507 - val_loss: 38.2744
Epoch 26/50
60000/60000 [=====] - 8s 134us/step - loss:
37.7855 - val_loss: 38.2082
Epoch 27/50
60000/60000 [=====] - 8s 141us/step - loss:
37.7191 - val_loss: 38.1145
Epoch 28/50
60000/60000 [=====] - 10s 160us/step - loss:
37.6795 - val_loss: 38.0433
Epoch 29/50
60000/60000 [=====] - 9s 146us/step - loss:
37.6100 - val_loss: 38.1676
Epoch 30/50
60000/60000 [=====] - 9s 142us/step - loss:
37.5586 - val_loss: 37.9974
Epoch 31/50
60000/60000 [=====] - 9s 144us/step - loss:
37.5083 - val_loss: 38.0194
Epoch 32/50
60000/60000 [=====] - 9s 148us/step - loss:
37.4622 - val_loss: 37.9582
Epoch 33/50

60000/60000 [=====] - 9s 143us/step - loss:
37.4099 - val_loss: 37.9270
Epoch 34/50
60000/60000 [=====] - 8s 140us/step - loss:
37.3614 - val_loss: 37.9836
Epoch 35/50
60000/60000 [=====] - 9s 142us/step - loss:
37.3170 - val_loss: 37.9364
Epoch 36/50
60000/60000 [=====] - 8s 138us/step - loss:
37.2814 - val_loss: 37.8243
Epoch 37/50
60000/60000 [=====] - 8s 139us/step - loss:
37.2315 - val_loss: 37.8536
Epoch 38/50
60000/60000 [=====] - 9s 143us/step - loss:
37.2053 - val_loss: 37.9338
Epoch 39/50
60000/60000 [=====] - 8s 141us/step - loss:
37.1470 - val_loss: 37.9156
Epoch 40/50
60000/60000 [=====] - 8s 137us/step - loss:
37.0987 - val_loss: 37.7732
Epoch 41/50
60000/60000 [=====] - 9s 147us/step - loss:
37.0719 - val_loss: 37.8128
Epoch 42/50
60000/60000 [=====] - 9s 143us/step - loss:
37.0450 - val_loss: 37.7106
Epoch 43/50
60000/60000 [=====] - 9s 143us/step - loss:
37.0032 - val_loss: 37.7626
Epoch 44/50
60000/60000 [=====] - 8s 138us/step - loss:
36.9560 - val_loss: 37.7558
Epoch 45/50
60000/60000 [=====] - 8s 140us/step - loss:
36.9417 - val_loss: 37.7912
Epoch 46/50
60000/60000 [=====] - 9s 147us/step - loss:
36.9385 - val_loss: 37.7411
Epoch 47/50
60000/60000 [=====] - 9s 145us/step - loss:
36.8763 - val_loss: 37.7623
Epoch 48/50
60000/60000 [=====] - 9s 142us/step - loss:
36.8652 - val_loss: 37.7362
Epoch 49/50
60000/60000 [=====] - 9s 144us/step - loss:
36.8307 - val_loss: 37.7160
Epoch 50/50


```
60000/60000 [=====] - 9s 144us/step - loss:
36.7962 - val_loss: 37.6770
```

```
<keras.callbacks.callbacks.History at 0x63fa5e048>
```

```
[98] vae.save_weights('vae_mlp_mnist.h5')
```

```
[112] import os
import matplotlib.pyplot as plt
```

```
[128] def plot_results(models,
                    data,
                    batch_size=128,
                    model_name="vae_mnist"):
    """Plots labels and MNIST digits as a function of the 2D
    latent vector
    # Arguments
        models (tuple): encoder and decoder models
        data (tuple): test data and label
        batch_size (int): prediction batch size
        model_name (string): which model is using this function
    """

    encoder, decoder = models
    x_test, y_test = data
    os.makedirs(model_name, exist_ok=True)

    filename = os.path.join(model_name, "vae_mean.png")
    # display a 2D plot of the digit classes in the latent space
    z_mean, _, _ = encoder.predict(x_test,
                                   batch_size=batch_size)

    plt.figure(figsize=(10, 10))
    plt.scatter(z_mean[:, 0], z_mean[:, 1], c=y_test)
    plt.colorbar()
    plt.xlabel("z[0]")
    plt.ylabel("z[1]")
    plt.savefig(filename)
    plt.show()

    filename = os.path.join(model_name, "digits_over_latent.png")
    # display a 30x30 2D manifold of digits
    n = 30
    digit_size = 28
    figure = np.zeros((digit_size * n, digit_size * n))
    # linearly spaced coordinates corresponding to the 2D plot
    # of digit classes in the latent space
    grid_x = np.linspace(-2, 2, n)
    grid_y = np.linspace(-2, 2, n)[::-1]
```

```

for i, yi in enumerate(grid_y):
    for j, xi in enumerate(grid_x):
        z_sample = np.array([[xi, yi]])
        x_decoded = decoder.predict(z_sample)
        digit = x_decoded[0].reshape(digit_size, digit_size)
        figure[i * digit_size: (i + 1) * digit_size,
                j * digit_size: (j + 1) * digit_size] = digit

plt.figure(figsize=(10, 10))
start_range = digit_size // 2
end_range = (n - 1) * digit_size + start_range + 1
pixel_range = np.arange(start_range, end_range, digit_size)
sample_range_x = np.round(grid_x, 1)
sample_range_y = np.round(grid_y, 1)
plt.xticks(pixel_range, sample_range_x)
plt.yticks(pixel_range, sample_range_y)
plt.xlabel("z[0]")
plt.ylabel("z[1]")
plt.imshow(figure, cmap='Greys_r')
plt.savefig(filename)
plt.show()

```

```


[129] models = (encoder, decoder)
      data = (x_test, y_test)


```

```

[130] plot_results(models,
                data,
                batch_size=batch_size,
                model_name="vae_mlp")

```





```
[125] # from above -1,1 and -1,1 is pretty tight
      # -2,2 and -2,2 leaves the corners empty
      # 0 0 2
      # 0 6 2
      # 6532
      # 4 8 1
      # 49851
      # 7 9 1
```

```
[ ]
```